

KUSH PATEL

Data Analyst | BI & Analytics Engineer | Power BI · SQL · Python · Microsoft Fabric · dbt

Vancouver, BC | +1 (778) 637-1602 | dharma.patel552@gmail.com | linkedin.com/in/kush-patel-48885719b | github.com/KushPatel29 | kushpatel29.github.io

Open to remote and relocation across Canada · Legally eligible to work in Canada · Available immediately

SUMMARY

- Analytics engineer and BI analyst with 4+ years turning raw operational data — ERP, POS, WMS, CRM, REST feeds — into governed warehouse layers and the reports the business argues from. Sole analyst across a perishable-goods supply chain, owning everything from month-end GL reconciliation to lot-level traceability.
- Builds dimensional models on cloud warehouses (Microsoft Fabric, Azure Synapse, SQL Server) with Bronze/Silver/Gold medallion layers, conformed dimensions and SCD Type 1/2, orchestrated in Azure Data Factory and Fabric pipelines (Airflow in project work, not production).
- Delivers Power BI end to end — semantic model, DAX measure library, RLS/OLS, publish-ready reports — so Finance, Sales and Supply Chain self-serve from one definition instead of competing extracts.
- Treats data quality as engineering. Thirteen featured projects carry **8,887 automated tests** in CI, checking documented invariants and benchmark results. Fourteen public project repositories include reproducible synthetic data and implementation notes.

TECHNICAL SKILLS

Warehousing & Lakehouse: Microsoft Fabric (Lakehouse, Warehouse, OneLake, Direct Lake), Azure Synapse Analytics, SQL Server, PostgreSQL, DuckDB, Delta Lake, Snowflake (dual-target dbt profile)

ELT / Orchestration: Azure Data Factory, Fabric Data Pipelines & Dataflows Gen2, dbt Core, Apache Airflow, SSIS, SQL Server Agent, incremental and CDC-style loads, schema-drift contracts, quarantine-and-replay

BI & Semantic Layer: Power BI (DAX, Power Query, TMDL/PBIR, calculation groups, RLS/OLS, VertiPaq tuning), SSRS paginated reporting, Tableau, Streamlit, Plotly

Languages: SQL (T-SQL, window functions, CTEs), Python (pandas, NumPy, scikit-learn, Flask, FastAPI), PySpark, DAX, M

Modelling & Governance: Kimball star schemas, conformed dimensions, SCD Type 1/2, metric catalogues with owner and grain, data lineage, RLS/RBAC and column masking, de-identification (Safe Harbor, k-anonymity), cutover runbooks

Quality & Delivery: pytest, dbt tests, pandera contracts, referential-integrity and control-total checks, anomaly thresholds, reconciliation gates, Git, GitHub Actions, Docker

PROFESSIONAL EXPERIENCE

Two Rivers Specialty Meats

Dec 2023 – Jun 2026 · North Vancouver, BC

Data Analyst — Operations, Logistics & Enterprise Reporting

- Sole analyst for a perishable-goods distributor: integrated ERP, Square POS, workforce and REST API feeds into curated Bronze/Silver/Gold datasets serving ad hoc analysis, BI dashboards and operational reporting.
- Designed and built the enterprise data mart applying star schemas, conformed dimensions and slowly changing dimensions, producing reliable data products sliced by SKU, supplier, region, customer and channel.
- Owned Power BI end to end — semantic model design, a standardised DAX measure library, drill-through and publish-ready reports — giving Sales, Finance and Supply Chain self-service access without analyst dependency.
- Engineered cloud-native ELT in Microsoft Fabric and Azure Data Factory/Synapse using SQL and PySpark, with incremental loads, parameterised transforms and structured error handling.
- Built reconciliation controls comparing source against curated totals for GL/P&L and inventory valuation, cutting discrepancies an estimated ~45% and reporting errors ~30% by resolving issues at the feed level rather than patching reports (internal figures, directional rather than audited).
- Maintained data quality across the mart: automated schema, completeness, uniqueness and referential-integrity tests, anomaly thresholds, exception logging and pipeline alerting.
- Implemented FEFO and lot-level traceability for recall readiness, delivered customer intelligence reporting (inactivity signals, lost-account tracking, rep coverage), and built SKU-level demand forecasting deployed as Python web apps that reduced stock-outs and perishable waste.
- Owned the legacy SQL Server estate through migration: maintained SSIS packages and SSRS paginated reports while refactoring priority pipelines to Fabric notebooks, validating cutover with row counts, checksums and control totals so no reporting was disrupted.
- Partnered with Sales, Finance and Supply Chain to define governed KPIs (revenue, COGS, gross margin, OTIF, inventory turns, days on hand) with documented definitions and acceptance criteria, and secured them with RLS/RBAC and column masking as self-serve access widened.

Shivam Investments

Dec 2020 – Jul 2022 · India / Remote

Financial Analyst — Data & Reporting

- Gathered reporting requirements from portfolio and risk stakeholders and delivered Power BI dashboards with drill-through and auditable measure definitions.
- Automated recurring analysis with Python (pandas) and SQL views feeding Power BI, replacing manual spreadsheet consolidation and cutting manual effort an estimated ~40%.
- Performed GL/P&L reconciliations across sources and investigated discrepancies to maintain finance-grade accuracy.

Retail Analytics Platform — Python, Flask, DuckDB, Plotly, Playwright, Docker · 1,412 tests
github.com/KushPatel29/wholesale-analytics-platform · live: kushpatel29.github.io/wholesale-analytics-platform
Sixteen pages that printed three different HHI values at once, because every page did its own arithmetic. One catalogue now governs 62 metric definitions — 47 implemented, 15 explicitly unavailable without a named source system — each carrying formula, grain, source table, owner page and gross-or-net basis, pinned by a hand-computed test.

Ask Your Data — grounded text-to-SQL — Python, DuckDB, ONNX, Claude API, Streamlit · 1,099 tests
github.com/KushPatel29/ask-your-data · live: ask-your-data-kp.streamlit.app
Ask 71 tables across 11 business domains a question, by text or voice, and audit the SQL behind every number. A deterministic compiler profiles the warehouse and binds language to it with no API key — 46 right, 0 wrong, 12 refused on a 58-question contract, because refusing is the feature. Model-written prose is checked back against the returned rows, so a sentence cannot quote a number the query never produced.

Supply Chain Control Tower — Microsoft Fabric, PySpark, Delta Lake, Power BI · 659 tests
github.com/KushPatel29/supply-chain-control-tower
A lakehouse medallion pipeline with a three-tier defence: versioned data contracts kill structural drift before Bronze, quarantine-and-replay isolates toxic rows, and a 20-check data-quality gate guards Gold — a critical failure halts publication. Adds global sourcing and King's-formula reorder points. Dynamic RLS and OLS verified by DAX impersonation against a pandas cross-check.

GL / P&L Reconciliation & Cloud Chargeback — Python, SQL, Power BI, Tableau · 558 tests
github.com/KushPatel29/gl-reconciliation-dashboard
Turns month-end into an audit trail instead of a manual variance hunt: 95.7% match rate, 230 exceptions and \$823K impact, each classified missing, timing, amount or duplicate, with ten accounts flagged out of tolerance and a documented SOX control. A FinOps second act maps FOCUS-shaped cloud billing through the same engine unmodified, with allocation coverage tying to the penny. Ships a generated Tableau workbook verified against the Python results.

Health System Decision Support — Python, Power BI, DAX, SPC · 604 tests
github.com/KushPatel29/healthcare-claims-analytics
CIHI-DAD-shaped activity with cost per weighted case, ALC bed equivalents and statistical process control using Laney's correction — because the readmission metric everyone reports is 4.6× overdispersed and a naive chart fires 41 false signals across 19 of 24 months. The US side prices \$3.8M of open AR at the ~\$1.7M it will actually collect. Fully synthetic; no PHI.

Legacy-to-Fabric Migration — SQL Server, SSIS, SSRS, T-SQL, Fabric · 203 tests
github.com/KushPatel29/legacy-to-fabric-migration
A cutover nobody has to take on faith: legacy SSIS/SSRS and the Fabric replacement run side by side, with row counts, checksums and control totals compared automatically and a gate that decides when the old system can be retired.

EDUCATION

Northeastern University — Master of Professional Studies in Analytics (Applied Machine Intelligence), GPA 3.76/4.0 · Vancouver, BC	Sep 2022 – Mar 2024
Gujarat Technological University — Bachelor of Engineering, Computer Science · India	Aug 2017 – Aug 2021

CERTIFICATIONS

Microsoft Certified: Fabric Analytics Engineer Associate (DP-600) — in progress · Google Data Analytics Professional Certificate · UC Davis SQL for Data Science · IBM Python for Data Science, AI & Development · Databricks Lakehouse Fundamentals · Power BI Data Modelling with DAX · Northeastern Data Visualization